

# Condor Administration

Alan De Smet

Computer Sciences Department  
University of Wisconsin-Madison

[condor-admin@cs.wisc.edu](mailto:condor-admin@cs.wisc.edu)

<http://www.cs.wisc.edu/condor>



# Outline

- Condor Daemons
  - Job Startup
- Configuration Files
- Policy Expressions
  - Startd (Machine)
  - Negotiator
- Priorities
- Security
- Administration
- Installation
  - "Full Installation"
- Other Sources

# Condor Daemons

# Condor Daemons

- > `condor_master` - controls everything else
- > `condor_startd` - executing jobs
  - `condor_starter` - helper for starting jobs
- > `condor_schedd` - submitting jobs
  - `condor_shadow` - submit-side helper

# Condor Daemons

- > `condor_collector` - Collects system information; only on Central Manager
- > `condor_negotiator` - Assigns jobs to machines; only on Central Manager

# Condor Daemons

- > You only have to run the daemons for the services you want to provide
- > **DAEMON\_LIST** is a comma separated list of daemons to start
  - **DAEMON\_LIST=MASTER, SCHEDD, STARTD**

# condor\_master

- Starts up all other Condor daemons
- If a daemon exits unexpectedly, restarts daemon and emails administrator
- If a daemon binary is updated (timestamp changed), restarts the daemon

# condor\_master

- Provides access to many remote administration commands:
  - `condor_reconfig`, `condor_restart`, `condor_off`, `condor_on`, etc.
- Default server for many other commands:
  - `condor_config_val`, etc.





# condor\_master

- Periodically runs `condor_preen` to clean up any files Condor might have left on the machine
  - Backup behavior, the rest of the daemons clean up after themselves, as well

# condor\_startd

- Represents a machine to the Condor pool
- Should be run on any machine you want to run jobs
- Enforces the wishes of the machine owner (the owner's "policy")

# condor\_startd

- > Starts, stops, suspends jobs
- > Spawns the appropriate `condor_starter`, depending on the type of job
- > Provides other administrative commands (for example, `condor_vacate`)

# condor\_starter

- Spawned by the `condor_startd` to handle all the details of starting and managing the job
  - Transfer job's binary to execute machine
  - Send back exit status
  - Etc.

# condor\_starter

- On multi-processor machines, you get one `condor_starter` per CPU
  - Actually one per running job
  - Can configure to run more (or less) jobs than CPUs
- For PVM jobs, the starter also spawns a PVM daemon (`condor_pvmd`)

# condor\_schedd

- > Represents jobs to the Condor pool
- > Maintains persistent queue of jobs
  - Queue is not strictly FIFO (priority based)
  - Each machine running condor\_schedd maintains its own queue
- > Should be run on any machine you want to submit jobs from



# condor\_schedd

- Responsible for contacting available machines and spawning waiting jobs
  - When told to by condor\_negotiator
- Services most user commands:
  - `condor_submit`, `condor_rm`,  
`condor_q`

# condor\_shadow

- Represents job on the submit machine
- Services requests from standard universe jobs for remote system calls
  - including all file I/O
- Makes decisions on behalf of the job
  - for example: where to store the checkpoint file



# condor\_shadow Impact

- One condor\_shadow running on submit machine for each actively running Condor job
- Minimal load on submit machine
  - Usually blocked waiting for requests from the job or doing I/O
  - Relatively small memory footprint

# Limiting condor\_shadow

- Still, you can limit the impact of the shadows on a given submit machine:
  - They can be started by Condor with a “nice-level” that you configure (`SHADOW_NICE_INCREMENT`)
  - Can limit total number of shadows running on a machine (`MAX_JOBS_RUNNING`)



# condor\_collector

- Collects information from all other Condor daemons in the pool
- Each daemon sends a periodic update called a ClassAd to the collector
- Services queries for information:
  - Queries from other Condor daemons
  - Queries from users (condor\_status)

# condor\_negotiator

- Performs matchmaking in Condor
  - Pulls list of available machines and job queues from `condor_collector`
  - Matches jobs with available machines
  - Both the job and the machine must satisfy each other's requirements (2-way matching)
- Handles user priorities



# Central Manager

- > The Central Manager is the machine running the collector and negotiator

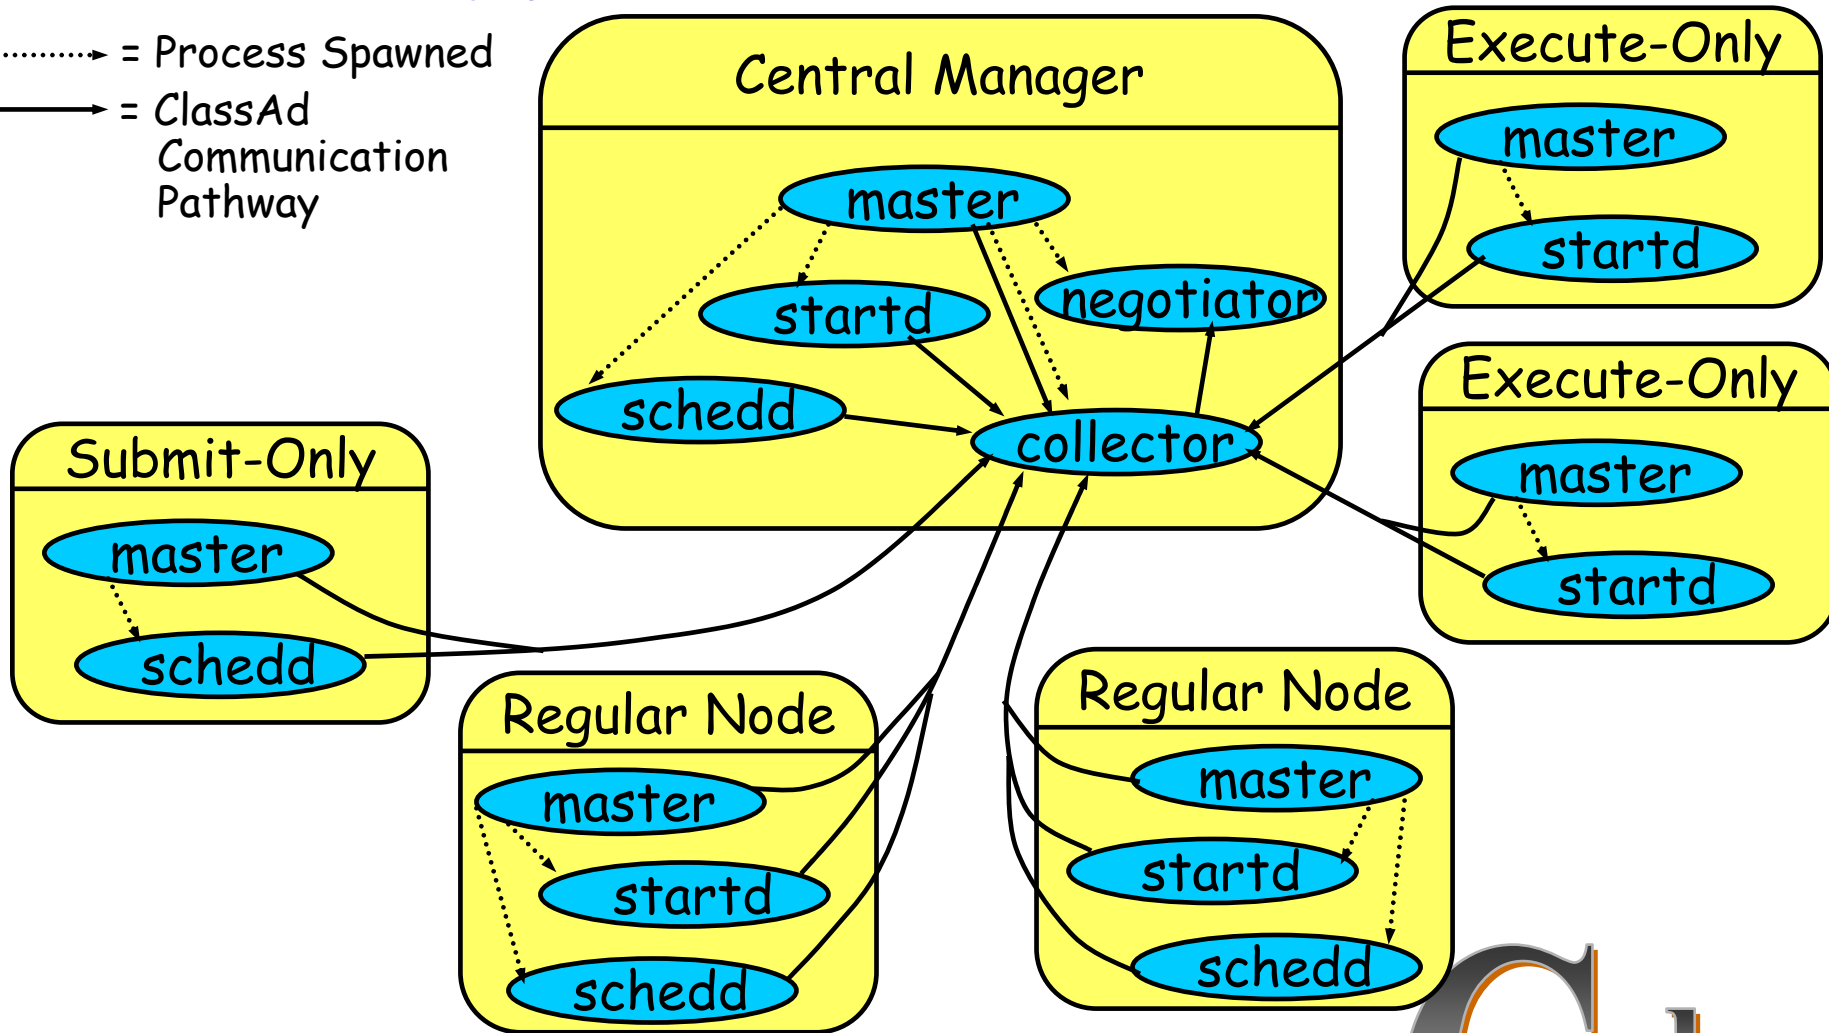
```
DAEMON_LIST = MASTER,  
COLLECTOR, NEGOTIATOR
```

- > Defines a Condor pool.

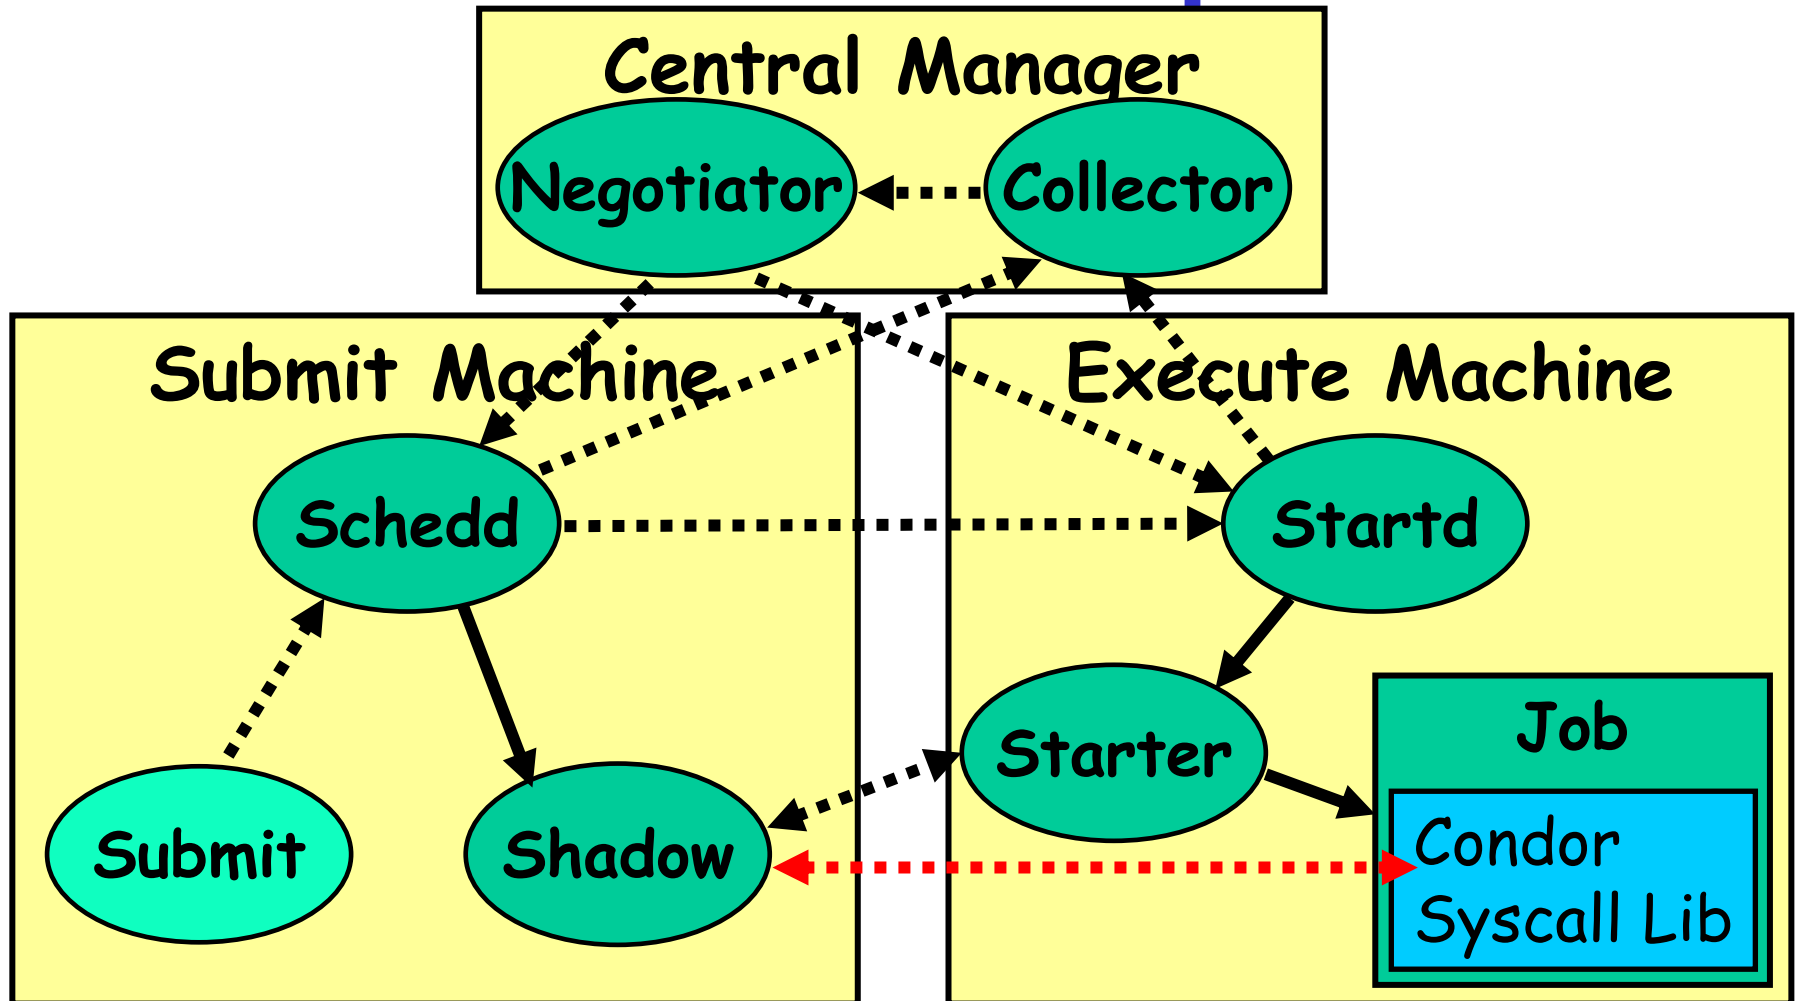
```
CONDOR_HOST =  
centralmanager.example.com
```

# Typical Condor Pool

.....> = Process Spawned  
——> = ClassAd Communication Pathway



# Job Startup



# Configuration Files



# Configuration Files

- Multiple files concatenated
  - Definitions in later files overwrite previous definitions
- Order of files:
  - Global config file
  - Local config files, shared config files
  - Global and Local Root config file

# Global Config File

- Found either in file pointed to with the `CONDOR_CONFIG` environment variable,  
`/etc/condor/condor_config`, or  
`~condor/condor_config`
- Most settings can be in this file
- Only works as a global file if it is on a shared file system

# Other Shared Files

- `LOCAL_CONFIG_FILE` macro
  - Comma separated, processed in order
- You can configure a number of other shared config files:
  - Organize common settings (for example, all policy expressions)
  - platform-specific config files

# Local Config File

- `LOCAL_CONFIG_FILE` macro (again)
  - Usually uses `$(HOSTNAME)`
- Machine-specific settings
  - local policy settings for a given owner
  - different daemons to run (for example, on the Central Manager!)

# Local Config File

- > Can be on local disk of each machine

`/var/adm/condor/condor_config.local`

- > Can be in a shared directory

`/shared/condor/condor_config.$(HOSTNAME)`

`/shared/condor/hosts/$(HOSTNAME) /  
condor_config.local`

# Root Config File (optional)

- Always processed last
- Allows root to specify settings which cannot be changed by other users
  - For example, the path to Condor daemons
- Useful if daemons are started as root but someone else has write access to config files

# Root Config File (optional)

- `/etc/condor/condor_config.root` or  
`~condor/condor_config.root`
- Then loads any files specified in  
`ROOT_CONFIG_FILE_LOCAL`

# Configuration File Syntax

- # at start of line is a comment
  - not allowed in names, confuses Condor.
- \ at the end of line is a line-continuation
  - Both lines are treated as one big entry
  - Works in comments!



# Configuration File Macros

- > Macros have the form:
  - `Attribute_Name = value`
    - Names are case insensitive
    - Values are case sensitive
- > You reference other macros with:
  - `A = $(B)`
- > Can create additional macros for organizational purposes

# Configuration File Macros

- > Can append to macros:

`A=abc`

`A=$(A),def`

- > Don't let macros recursively define each other!

`A=$(B)`

`B=$(A)`

# Configuration File Macros

- Later macros in a file overwrite earlier ones

- B will evaluate to 2:

A=1

B=\$ ( A )

A=2

# ClassAds

- Set of key-value pairs
- Can be matched against each other
  - Requirements and Rank
- This is old ClassAds
  - New, more expressive ClassAds exist
    - Not yet used in Condor

# ClassAd Expressions

- Some configuration file macros specify expressions for the Machine's ClassAd
  - Notably START, RANK, SUSPEND, CONTINUE, PREEMPT, KILL
- Can contain a mixture of macros and ClassAd references
- Notable: UNDEFINED, ERROR

# ClassAd Expressions

- > +, -, \*, /, <, <=, >, >=, ==, !=, &&, and || all work as expected
- > TRUE==1 and FALSE==0 (guaranteed)

# Macros and Expressions Gotcha

- > These are simple replacement macros
- > Put parentheses around expressions

**TEN=5+5**

**HUNDRED=\$ ( TEN ) \* \$ ( TEN )**

- HUNDRED becomes 5+5\*5+5 or 35!

**TEN= ( 5+5 )**

**HUNDRED= ( \$ ( TEN ) \* \$ ( TEN ) )**

- $((5+5)*(5+5)) = 100$

# ClassAd Expressions: UNDEFINED and ERROR

- > Special values
- > Passed through most operators
  - Anything == UNDEFINED is UNDEFINED
- > && and || eliminate if possible.
  - UNDEFINED && FALSE is FALSE
  - UNDEFINED && TRUE is UNDEFINED



# ClassAd Expressions: `=?=` and `!=`

- `=?=` and `!=` are similar to `==` and `!=`
- `=?=` tests if operands have the same type and the same value.
  - `10 == UNDEFINED -> UNDEFINED`
  - `UNDEFINED == UNDEFINED -> UNDEFINED`
  - `10 ==?= UNDEFINED -> FALSE`
  - `UNDEFINED ==?= UNDEFINED -> TRUE`
- `!=` inverts `==`

# ClassAd Expressions

- Further information: Section 4.1, "Condor's ClassAd Mechanism," in the Condor Manual.

# Policy Expressions

# Policy Expressions

- Allow machine owners to specify job priorities, restrict access, and implement local policies

# Policy Expressions

- Specified in `condor_config`
- Policy evaluates both a machine ClassAd and a job ClassAd together
  - Policy can reference items in either ClassAd (See manual for list)
- Can reference `condor_config` macros: `$(MACRONAME)`

# Machine (Startd) Policy Expression Summary

- **START** - When is this machine willing to start a job
  - Typically used to restrict access when the machine is being used directly
- **RANK** - Job preferences

# Machine (Startd) Policy Expression Summary

- **SUSPEND** - When to suspend a job
- **CONTINUE** - When to continue a suspended job
- **PREEMPT** - When to nicely stop running a job
- **KILL** - When to immediately kill a preempting job

# START

- START is the primary policy
- When FALSE the machine enters the Owner state and will not run jobs
- Acts as the Requirements expression for the machine, the job must satisfy START
  - Can reference job ClassAd values including Owner and ImageSize



# RANK

- Indicates which jobs a machine prefers
  - Jobs can also specify a rank
- Floating point number
  - Larger numbers are higher ranked
  - Typically evaluate attributes in the Job ClassAd
  - Typically use + instead of &&

# RANK

- Often used to give priority to owner of a particular group of machines
- Claimed machines still advertise looking for higher ranked job to preempt the current job

# SUSPEND and CONTINUE

- > When SUSPEND becomes true, the job is suspended
- > When CONTINUE becomes true a suspended job is released

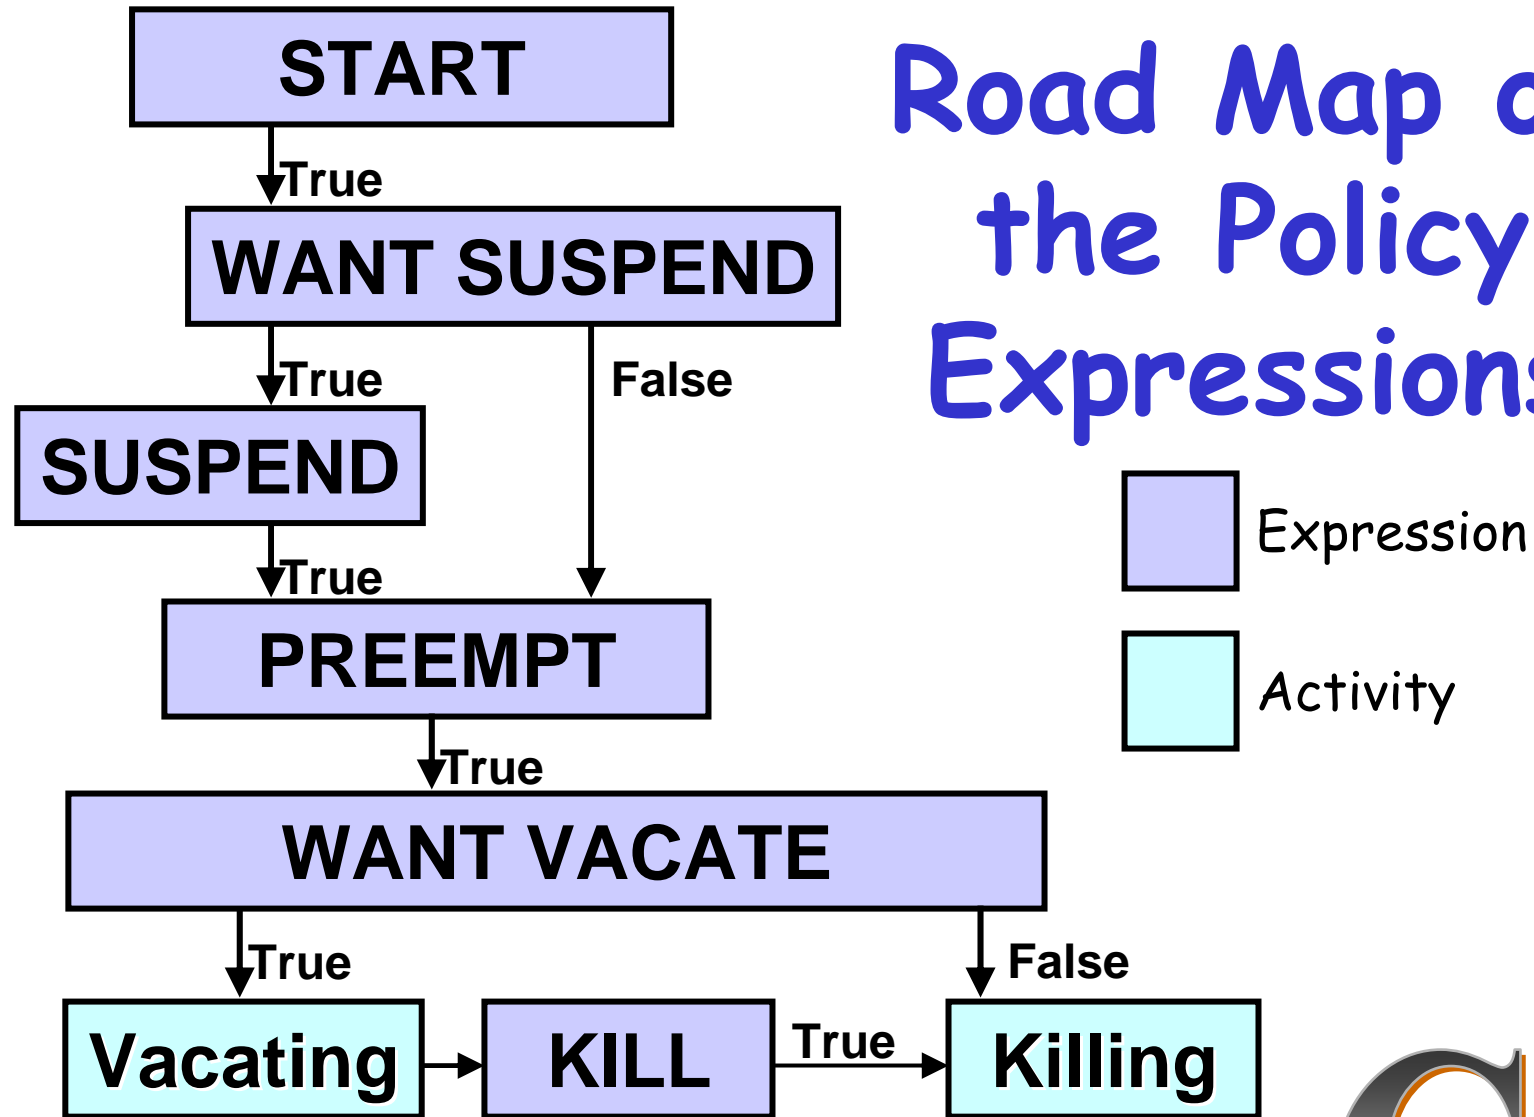
# PREEMPT and KILL

- When PREEMPT becomes true, the job will be politely shut down
  - Vanilla universe jobs get SIGTERM
  - Standard universe jobs checkpoint
- When KILL becomes true, the job is SIGKILL
  - Checkpointing is aborted if started

# WANT\_SUSPEND and WANT\_VACATE

- > Typically leave both to TRUE
- > WANT\_SUSPEND - If false, skip SUSPEND test, jump to PREEMPT
- > WANT\_VACATE
  - If true, gives job time to vacate cleanly (until KILL becomes true)
  - If false, job is immediately killed (KILL is ignored)

# Road Map of the Policy Expressions



# Minimal Settings

> Always runs jobs

**START** = True

**RANK** =

**SUSPEND** = False


**CONTINUE** = True

**PREEMPT** = False

**KILL** = False

# Policy Configuration

(Boss Fat Cat)



- > I am adding nodes to the Cluster... *but the Chemistry Department has priority on these nodes*



# New Settings for the Chemistry nodes

## > Prefer Chemistry jobs

**START** = True

**RANK** = *Department* == "Chemistry"

**SUSPEND** = False

**CONTINUE** = True

**PREEMPT** = False

**KILL** = False

# Submit file with Custom Attribute

- > Prefix an entry with "+" to add to job ClassAd

`Executable = charm-run`

`Universe = standard`

`+Department = Chemistry`  
`queue`

# What if "Department" not specified?

**START** = True

**RANK** = *Department != UNDEFINED  
&& Department == "Chemistry"*

**SUSPEND** = False

**CONTINUE** = True

**PREEMPT** = False

**KILL** = False

# More Complex RANK

- > Give the machine's owners (adesmet and roy) highest priority, followed by the Chemistry department, followed by the Physics department, followed by everyone else.

# More Complex RANK

```
IsOwner = (Owner == "adesmet" ||  
            Owner == "roy")
```


```
IsChem =(Department != UNDEFINED  
            && Department == "Chemistry")
```

```
IsPhys =(Department != UNDEFINED  
            && Department == "Physics")
```

```
RANK = $(IsOwner)*20 +  
         $(IsChem)*10 + $(IsPhys)
```

# Policy Configuration

(Boss Fat Cat)



- > Cluster is okay, but... *Condor can only use the desktops when they would otherwise be idle*

# Defining Idle

- One possible definition:
  - No keyboard or mouse activity for 5 minutes
  - Load average below 0.3

# Desktops should

- > **START** jobs when the machine becomes idle
- > **SUSPEND** jobs as soon as activity is detected
- > **PREEMPT** jobs if the activity continues for 5 minutes or more
- > **KILL** jobs if they take more than 5 minutes to preempt



# Macros in the Config File

**NonCondorLoadAvg** = (LoadAvg - CondorLoadAvg)

**HighLoad** = 0.5

**BgndLoad** = 0.3

**CPU\_Busy** = (\$ (NonCondorLoadAvg) >=  
\$ (HighLoad) )

**CPU\_Idle** = (\$ (NonCondorLoadAvg) <=  
\$ (BgndLoad) )

**KeyboardBusy** = (KeyboardIdle < 10)

**MachineBusy** = (\$ (CPU\_Busy) ||  
\$ (KeyboardBusy) )

**ActivityTimer** = \  
(CurrentTime - EnteredCurrentActivity)

# Desktop Machine Policy

**START** = \$(CPU\_Idle) && KeyboardIdle > 300

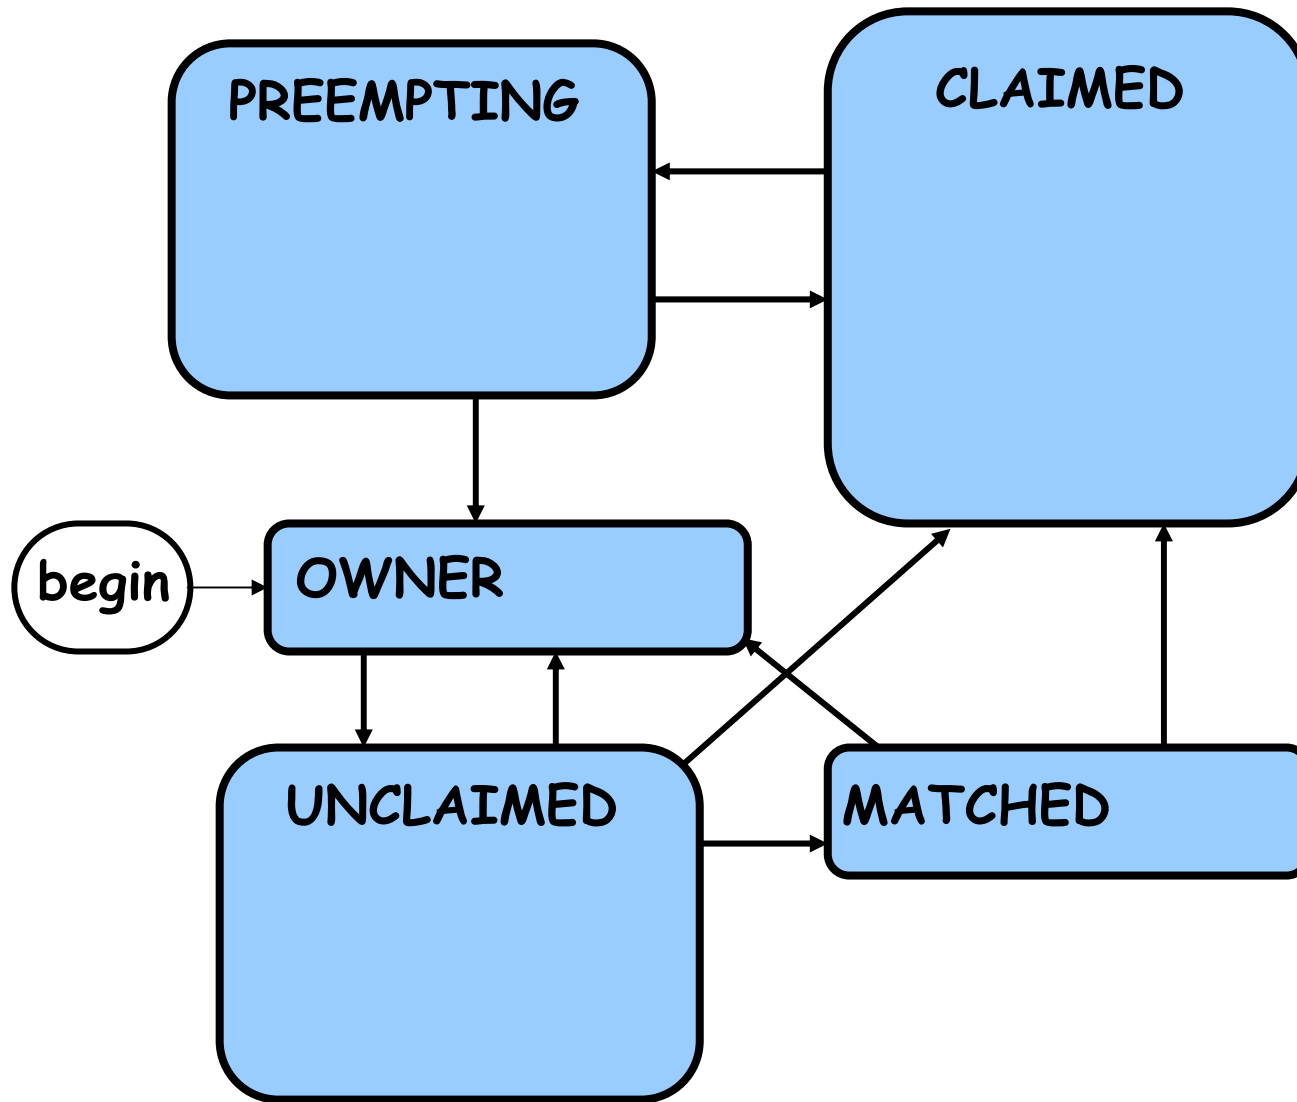
**SUSPEND** = \$(MachineBusy)

**CONTINUE** = \$(CPU\_Idle) && KeyboardIdle > 120

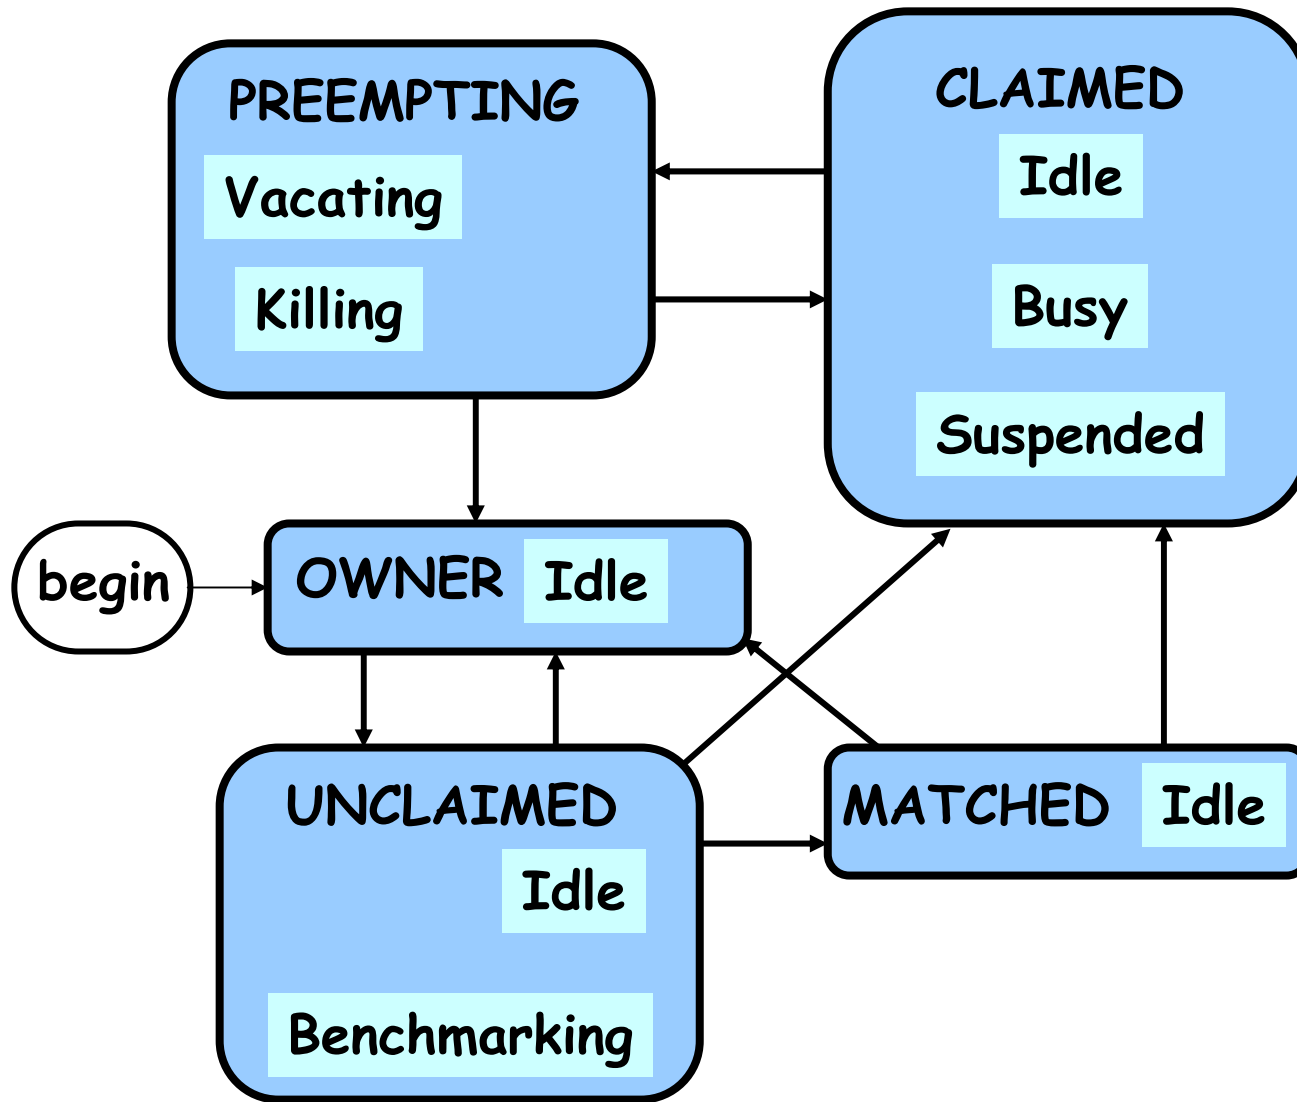
**PREEMPT** = (Activity == "Suspended") && \  
\$(ActivityTimer) > 300

**KILL** = \$(ActivityTimer) > 300

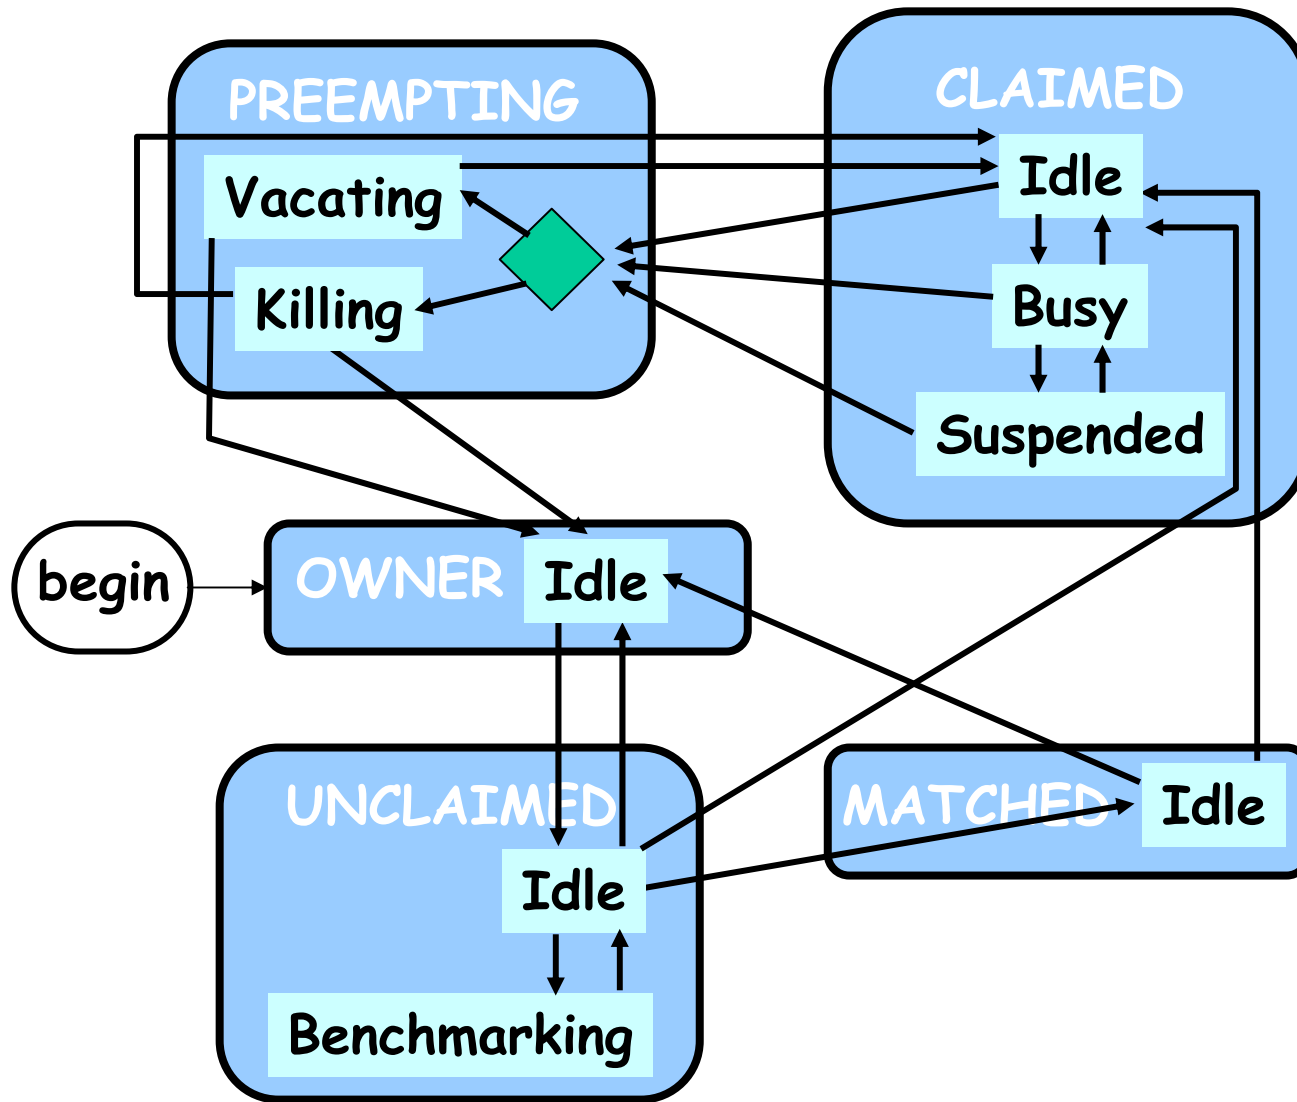
# Machine States



# Machine Activities



# Machine Activities



See the manual for the gory details (Section 3.6: Configuring the Startd Policy)

# Custom Machine Attributes

- Can add attributes to a machine's ClassAd, typically done in the local config file

**INSTRUCTIONAL=TRUE**

**NETWORK\_SPEED=100**

**STARTD\_EXPRS=INSTRUCTIONAL,  
NETWORK\_SPEED**

# Custom Machine Attributes

- > Jobs can now specify Rank and Requirements using new attributes:

```
Requirements =  
    ( INSTRUCTIONAL=?=UNDEFINED ||  
      INSTRUCTIONAL==FALSE )
```

```
Rank = NETWORK_SPEED !=  
      UNDEFINED && NETWORK_SPEED
```

# Policy Review

- Users submitting jobs can specify Requirements and Rank expressions
- Administrators can specify Startd policy expressions individually for each machine
- Custom attributes easily added
- You can enforce almost any policy!



# Further Machine Policy Information

- > For further information, see section 3.6 "Startd Policy Configuration" in the Condor manual
- > condor-users mailing list  
<http://www.cs.wisc.edu/condor/mail-lists/>
- > [condor-admin@cs.wisc.edu](mailto:condor-admin@cs.wisc.edu)

# Priorities

# Job Priority

- Set with `condor_prio`
- Range from -20 to 20
  - Larger numbers are higher priority"
- Only impacts order between jobs for a single user

# User Priority

- Determines allocation of machines to waiting users
- View with `condor_userprio`
- Inversely related to machines allocated
  - A user with priority of 10 will be able to claim twice as many machines as a user with priority 20

# User Priority

- Effective User Priority is determined by multiplying two factors
  - Real Priority
  - Priority Factor

# Real Priority

- Based on actual usage
- Defaults to 0.5
- Approaches actual number of machines used over time
  - Configuration setting  
`PRIORITY_HALFLIFE`

# Priority Factor

- Assigned by administrator
  - Set with `condor_userprio`
- Defaults to 1 (`DEFAULT_PRIO_FACTOR`)
- Nice users default to 1,000,000 (`NICE_USER_PRIO_FACTOR`)
  - Used for true bottom feeding jobs
  - Add "`nice_user=true`" to your submit file

# Negotiator Policy Expressions

- **PREEMPTION\_REQUIREMENTS** and **PREEMPTION\_RANK**
- Evaluated when `condor_negotiator` considers replacing a lower priority job with a higher priority job
- Completely unrelated to the **PREEMPT** expression



# PREEMPTION\_REQUIREMENTS

- If false will not preempt machine
  - Typically used to avoid pool thrashing

```
PREEMPTION_REQUIREMENTS = \
```

```
$(StateTimer) > (1 * $(HOUR)) \
```

```
&& RemoteUserPrio > SubmittorPrio * 1.2
```

- Only replace jobs running for at least one hour and 20% lower priority

# PREEMPTION\_RANK

- Picks which already claimed machine to reclaim

`PREEMPTION_RANK = \`

`(RemoteUserPrio * 1000000) \`

`- ImageSize`

- Strongly prefers preempting jobs with a large (bad) priority and a small image size

# Security

# Host/IP Address Security

- The basic security model in Condor
  - Stronger security available (Encrypted communications, cryptographic authentication)
- Can configure each machine in your pool to allow or deny certain actions from different groups of machines

# Security Levels

- READ access - querying information
  - `condor_status`, `condor_q`, etc
- WRITE access - updating information
  - Does *not* include READ access!
  - `condor_submit`, adding nodes to a pool, etc

# Security Levels

## > ADMINISTRATOR access

- `condor_on`, `condor_off`,  
`condor_reconfig`, `condor_restart`,  
etc.

## > OWNER access

- Things a machine owner can do (notably  
`condor_vacate`)

# Setting Up Security

- List what hosts are allowed or denied to perform each action
  - If you list allowed hosts, everything else is denied
  - If you list denied hosts, everything else is allowed
  - If you list both, only allow hosts that are listed in "allow" but not in "deny"

# Specifying Hosts

- There are many possibilities for specifying which hosts are allowed or denied:
  - Host names, domain names
  - IP addresses, subnets



# Wildcards

- > '\*' can be used anywhere (once) in a host name
  - for example, "infn-corsi\*.corsi.infn.it"
- > '\*' can be used at the end of any IP address
  - for example "128.105.101.\*" or "128.105.\*"

# Setting up Host/IP Address Security

- Can define values that effect all daemons:
  - `HOSTALLOW_WRITE`, `HOSTDENY_READ`, `HOSTALLOW_ADMINISTRATOR`, etc.
- Can define daemon-specific settings:
  - `HOSTALLOW_READ_SCHEDD`, `HOSTDENY_WRITE_COLLECTOR`, etc.

# Example Security Settings

```
HOSTALLOW_WRITE = *.infn.it
```

```
HOSTALLOW_ADMINISTRATOR= infn-corsi1*,\  
$(CONDOR_HOST), axpb07.bo.infn.it, \  
$(FULL_HOSTNAME)
```

```
HOSTDENY_ADMINISTRATOR = infn-corsi15
```

```
HOSTDENY_READ = *.gov, *.mil
```

```
HOSTDENY_ADMINISTRATOR_NEGOTIATOR = *
```



# Default Security Settings

```
HOSTALLOW_ADMINISTRATOR =  
$(CONDOR_HOST)
```

```
HOSTALLOW_OWNER = $(FULL_HOSTNAME),  
$(HOSTALLOW_ADMINISTRATOR)
```

```
HOSTALLOW_READ = *
```

```
HOSTALLOW_WRITE = *
```

> Make write restrictive

```
HOSTALLOW_WRITE=*.site.uk
```

# Advanced Security Features

- > **AUTHENTICATION** - Who is allowed
- > **ENCRYPTION** - Private communications, requires **AUTHENTICATION**.
- > **INTEGRITY** - Checksums

# Security Features

- Features individually set as **REQUIRED, PREFERRED, OPTIONAL, or NEVER**
- Can set default and for each level (**READ, WRITE, etc**)
- All default to **OPTIONAL**
- Leave **NEGOTIATION** at **OPTIONAL**

# Authentication Complexity

- Authentication comes at a price: complexity
- Authentication between machines requires an authentication system
- Condor supports several existing authentication systems
  - We don't want to create yet another one

# AUTHENTICATION\_METHODS

> Authentication requires one or more methods:

- FS
- FS\_REMOTE
- GSI
- Kerberos
- NTSSPI
- CLAIMTOBE



# FS and FS\_REMOTE Filesystem Tests

- FS checks that the user can create a file owned by the user.
  - Only works on local machine
  - Assumes the filesystem is trustworthy
- FS\_REMOTE works remotely
  - Allows test file to be on NFS, AFS, or other shared file system

# GSI

## Globus Security Infrastructure

- Daemons and users have X.509 certs
- All Condor daemons in pool can share one certificate
- Map file maps from X.509 distinguished names to identities.

# Kerberos and NTSSPI

- > Kerberos
  - Complex to set up
  - If you are already using, easy to add to Condor
- > NTSSPI - Windows NT
  - Only works on Windows

# CLAIMTOBE

- Trust any claims about user identity
  - If used, encryption's secret password passed in clear!
  - Use with care

# Additional Security Levels

## > CONFIG

- Dynamically change config settings

## > IMMEDIATE\_FAMILY

- Daemon to daemon communications

## > NEGOTIATOR

- `condor_negotiator` to other daemons

# ALLOW and DENY

- When authentication is enabled you can filter based on user identifier
- Use ALLOW and DENY instead of HOSTALLOW and HOSTDENY
- Can specify hostnames and IPs as before

# Specifying User Identities

- > username@site.example.com/hostname
- > Can use \* wildcard
- > Hostname can be hostname or IP address with optional netmask
  - 192.168.12.1/255.255.192.0
  - 192.168.12.1/18

# Example Filters

- > Allow anyone from wisc.edu:

```
ALLOW_READ=*@wisc.edu/*.wisc.edu
```

- > Allow any authorized local user:

```
ALLOW_READ=*/*.wisc.edu
```

- > Allow specific user/machine

```
ALLOW_NEGOTIATOR=
```

```
daemon@wisc.edu/condor.wisc.edu
```



# Example Advanced Security Configuration

- Enable authentication, encryption, and integrity
- Use GSI authentication for between machine connections
- Use GSI or FS authentication on a single machine

# Example Advanced Security Configuration

# Turn on all security:

SEC\_DEFAULT\_AUTHENTICATION=REQUIRED

SEC\_DEFAULT\_ENCRYPTION=REQUIRED

SEC\_DEFAULT\_INTEGRITY=REQUIRED

# Example Advanced Security Configuration

```
# Require authentication
```

```
SEC_DEFAULT_AUTHENTICATION_METHODS  
= FS, GSI
```

# Example Advanced Security Configuration

**ALLOW\_READ = \***

**ALLOW\_WRITE= \*@wisc.edu/\*.wisc.edu**

**DENY\_WRITE = abuser@\*.wisc.edu/\***

**ALLOW\_ADMINISTRATOR =**

**admin@wisc.edu/\*.wisc.edu,**

**\*@wisc.edu/\$(CONDOR\_HOST)**

# Example Advanced Security Configuration

```
ALLOW_CONFIG =  
$(ALLOW_ADMINISTRATOR)  
ALLOW_IMMEDIATE_FAMILY =  
daemon@wisc.edu/* .wisc.edu
```

# Example Advanced Security Configuration

```
ALLOW_OWNER =  
    $(ALLOW_ADMINISTRATOR) ,  
    $(FULL_HOSTNAME)
```

```
ALLOW_NEGOTIATOR =  
    daemon@wisc.edu/  
    $(CONDOR_HOST)
```

# Users without Certs

- Using FS authentication users can submit jobs and check the local queue
- `condor_status` won't work for normal users without an X.509 Cert
  - Requires READ access to `condor_collector`
- Can let anyone read any daemon!

# Allow Any User Read Access

# Using dreaded CLAIMTOBE

```
SEC_READ_AUTHENTICATION_METHODS =  
    FS, GSI, CLAIMTOBE
```



# Advanced Security Features

- For further details
  - Chapter 3.7, "Security in Condor" in the Condor Manual
  - [condor-admin@cs.wisc.edu](mailto:condor-admin@cs.wisc.edu)

# Administration

# condor\_config\_val

> Find current configuration values

```
% condor_config_val MASTER_LOG  
/var/condor/logs/MasterLog
```

# condor\_config\_val -v

> Can identify source

```
% condor_config_val -v CONDOR_HOST
```

```
CONDOR_HOST: condor.cs.wisc.edu
```

```
Defined in
```

```
`/etc/condor_config.hosts', line 6
```

# condor\_fetchlog

- > Retrieve logs remotely

`condor_fetchlog`

`beak.cs.wisc.edu Master`

# Querying daemons

## `condor_status`

- > Queries the collector for information about daemons in your pool
- > Defaults to finding `condor_startds`
- > `condor_status -schedd` summarizes all job queues
- > `condor_status -master` returns list of all `condor_masters`



# condor\_status

- > -long displays the full ClassAd
- > Specify a machine name to limit results to a single host

```
condor_status -l  
node4.cs.wisc.edu
```

## condor\_status -constraint

- > Only return ClassAds that match an expression you specify
- > Show me idle machines with 1GB or more memory
  - `condor_status -constraint 'Memory >= 1024 && Activity == "Idle"'`



# condor\_status -format

- Controls format of output
- Useful for writing scripts
- Uses C printf style formats
  - One field per argument

# condor\_status -format

> Census of systems in your pool:

```
% condor_status -format '%s'
Arch -format '%s\n' OpSys /
sort / uniq -c
```

```
797 INTEL LINUX
118 INTEL WINNT50
108 SUN4u SOLARIS28
6  SUN4x SOLARIS28
```

# Examining Queues

## `condor_q`

- View the job queue
- The “-long” option is useful to see the entire ClassAd for a given job
- supports `-constraint` and `-format`
- Can view job queues on remote machines with the “-name” option

# condor\_q -format

> Census of jobs per user

```
% condor_q -format '%8s' Owner  
-format '%s\n' Cmd | sort |  
uniq -c
```

```
64 adesmet /scratch/submit/a.out  
2 adesmet /home/bin/run_events  
4 smith /nfs/sim1/em2d3d  
4 smith /nfs/sim2/em2d3d
```

# condor\_q -analyze

- > condor\_q will try to figure out why the job isn't running
- > Good at determining that no machine matches the job Requirements expressions

# condor\_q -analyze

## > Typical results:

```
471216.000:  Run analysis summary.  Of 820 machines,  
    458 are rejected by your job's requirements  
    25 reject your job because of their own requirements  
    0 match, but are serving users with a better priority in the pool  
    4 match, but prefer another specific job despite its worse user-  
priority  
    6 match, but will not currently preempt their existing job  
    327 are available to run your job
```

# condor\_analyze

- Planned for 6.8
  - May be available as a contrib module for earlier releases
- Breaks down the job's requirements and suggests modifications

# condor\_analyze

## > (Heavily truncated output)

The Requirements expression for your job is:

```
( ( target.Arch == "SUN4u" ) && ( target.OpSys ==  
    "WINNT50" ) && [snip]
```

Condition	Machines	Suggestion
1 (target.Disk > 1000000000)	0	MODIFY TO 14223201
2 (target.Memory > 10000)	0	MODIFY TO 2047
3 (target.Arch == "SUN4u")	106	
4 (target.OpSys == "WINNT50")	110	MOD TO "SOLARIS28"

Conflicts: conditions: 3, 4



# Condor's Log Files

- Condor maintains one log file per daemon

# Condor's Log Files

- Can increase verbosity of logs on a per daemon basis
  - SHADOW\_DEBUG, SCHEDD\_DEBUG, and others
  - Space separated list

# Useful Debug Levels

- > `D_FULLDEBUG` dramatically increases information logged
- > `D_COMMAND` adds information about about commands received

```
SHADOW_DEBUG = \  
    D_FULLDEBUG D_COMMAND
```

# Condor's Log Files

- Log files are automatically rolled over when a size limit is reached
  - Defaults to 64000 bytes, you will probably want to increase.
  - Rolls over quickly with `D_FULLDEBUG`
  - `MAX_*_LOG`, one setting per daemon
    - `MAX_SHADOW_LOG`, `MAX_SCHEDD_LOG`, and others

# Condor's Log Files

- Many log files entries primarily useful to Condor developers
  - Especially if `D_FULLDEBUG` is on
  - Minor errors are often logged but corrected
  - Take them with a grain of salt
  - `condor-admin@cs.wisc.edu`

# Debugging Jobs: `condor_q`

- Examine the job with `condor_q`
  - especially `-long` and `-analyze`
  - Compare with `condor_status -long`

# Debugging Jobs: User Log

- Examine the job's user log

- Can find with:

```
condor_q -format '%s\n' UserLog 17.0
```

- Set with "log" in the submit file

- Contains the life history of the job

- Often contains details on problems

- Condor 6.6 includes improved messages

# Debugging Jobs: ShadowLog

- Examine ShadowLog on the submit machine
  - Note any machines the job tried to execute on
  - There is often an "ERROR" entry that can give a good indication of what failed



# Debugging Jobs: Matching Problems

- No shadowLog entries? Possible problem matching the job.
  - Examine scheddLog on the submit machine
  - Examine NegotiatorLog on the central manager

# Debugging Jobs: Local Problems

- > ShadowLog entries suggest an error but aren't specific?
  - Examine startLog and StarterLog on the execute machine

# Debugging Jobs: Reading Log Files

- Condor logs will note the job ID each entry is for
  - Useful if multiple jobs are being processed simultaneously
  - grepping for the job ID will make it easy to find relevant entries

# Debugging Jobs: What Next?

- If necessary add "D\_FULLDEBUG  
D\_COMMAND" to `DEBUG_DAEMONNAME`  
setting for additional log information
- Increase `MAX_DAEMONNAME_LOG` if  
logs are rolling over too quickly
- If all else fails, email us
  - `condor-admin@cs.wisc.edu`

# Installation

# Considerations for Installing a Condor Pool

- What machine should be your central manager?
- Does your pool have a shared file system?
- Where to install Condor binaries and configuration files?
- Where should you put each machine's local directories?
- Start the daemons as root or as some other user?

# What machine should be your central manager?

- The central manager is very important for the proper functioning of your pool
- If the central manager crashes, jobs that are currently matched will continue to run, but new jobs will not be matched

# Central Manager

- > Want assurances of high uptime or prompt reboots
- > A good network connection helps



# Does your pool have a shared file system?

- It is easier to run vanilla universe jobs if so, but one is not required
- Shared location for configuration files can ease administration of a pool
- AFS can work, but Condor does not yet manage AFS tokens

# Where to install binaries and configuration files?

- > Shared location for configuration files can ease administration of a pool
- > Binaries on a shared file system makes upgrading easier, but can be less stable if there are network problems
- > `condor_master` on the local disk is a good compromise

# Where should you put each machine's local directories?

- > You need a fair amount of disk space in the spool directory for each `condor_schedd` (holds job queue and binaries for each job submitted)
- > The execute directory is used by the `condor_starter` to hold the binary for any Condor job running on a machine

# Where should you put each machine's local directories?

- The log directory is used by all daemons
  - More space means more saved info

# Hostnames

- Any two machines that will be communicating must know each others names
- You can't have nameless machines

# Start the daemons as root or some other user?

- If possible, we recommend starting the daemons as root
  - Jobs run as the user that submitted them
    - More secure
    - Less confusion for users
  - Condor will try to run as the user "condor" whenever possible



# Running Daemons as Non-Root

- Condor will still work, users just have to take some extra steps to submit jobs
- Can have "personal Condor" installed - only you can submit jobs

# Basic Installation Procedure

- > 1. Decide what version and parts of Condor to install and download them
- > 2. Install the "release directory" - all the Condor binaries and libraries
- > 3. Setup the Central Manager
- > 4. (optional) Setup Condor on any other machines you wish to add to the pool
- > 5. Spawn the Condor daemons



# Condor Version Series

- We distribute two versions of Condor
  - Stable Series
  - Development Series

# Stable Series

- Heavily tested
- Recommended for general use
- 2nd number of version string is even (6.6.3)

# Development Series

- Latest features, not necessarily well-tested
- Not recommended unless you're willing to work with beta code or need new features
- 2nd number of version string is odd (6.7.0)

# Condor Versions

- > What am I running?
- > All daemons advertise a `CondorVersion` attribute in the `ClassAd` they publish
- > You can also view the version string by running `ident` on any Condor binary

# Condor Versions

- All parts of Condor on a single machine should run the same version!
- Machines in a pool can usually run different versions and communicate with each other
- Documentation will specify when a version is incompatible with older versions

# Downloading Condor

- Go to <http://www.cs.wisc.edu/condor/>
- Fill out the form and download the different pieces you need
  - Normally, you want the full stable release
- There are also “contrib” modules for non-standard parts of Condor
  - For example, the View Server

# Downloading Condor

- > Distributed as compressed "tar" files
- > Once you download, unpack them

# Install the Release Directory

- In the directory where you unpacked the tar file, you'll find a `release.tar` file with all the binaries and libraries
- `condor_configure` can help manage the installation



# condor\_configure

- > Handles installation and reconfiguration

```
condor_configure --install  
--install-dir=/nfs/opt/condor  
--local-dir=/var/condor  
--owner=condor
```

# Install the Release Directory

- > In a pool with a shared release directory, you should run `condor_configure` somewhere with write access to the shared directory
- > You need a *separate* release directory for *each platform*!

# Setup the Central Manager

- Central manager needs specific configuration to start the `condor_collector` and `condor_negotiator`
  - `condor_configure --type=manager`
    - or
  - `DAEMON_LIST = master, collector, negotiator`

# Setup Additional Machines

- If you have a shared file system, just run `condor_init` on any other machine you wish to add to your pool
  - Created local directories
- Without a shared file system, you must run `condor_configure` on each host

# Start the Condor daemons

- Run `condor_master` to start Condor
  - Remember to start as root if desired
- Start Condor on the central manager first
- Add Condor to your boot scripts?
  - We provide a "SysV-style" init script  
(`<release>/etc/examples/condor.boot`)

# Shared Release Directory

- Simplifies administration

# Shared Release Directory

- Unifies configuration files, simplifying changes
  - Same shared global config file for all machines
  - All local config files visible in one place
    - Can symlink local files for multiple machines to a single file

# Shared Release Directory

- Keep all of your binaries in one place
  - Prevents having different versions accidentally left on different machines
  - Easier to upgrade



# Condor-G Special Notes

- > Condor-G should work out of the box
- > Globus can push several limits, consider increasing:
  - /proc/sys/fs/file-max
  - /proc/sys/net/ipv4/ip\_local\_port\_range
  - Per process file descriptor limits

[http://www.cs.wisc.edu/condor/condorg/linux\\_scalability.html](http://www.cs.wisc.edu/condor/condorg/linux_scalability.html)

# “Full Installation” of `condor_compile`

- `condor_compile` re-links user jobs with Condor libraries to create “standard” universe jobs.
- By default, only works with certain commands (`gcc`, `g++`, `g77`, `cc`, `CC`, `f77`, `f90`, `ld`)
- With a “full-installation”, works with any command (notably, `make`)



# “Full Installation” of `condor_compile`

- Move real `ld` binary, the linker, to `ld.real`
  - Location of `ld` varies between systems, typically `/bin/ld`
- Install Condor's `ld` script in its place
- Transparently passes to `ld.real` by default; during `condor_compile` hooks in Condor libraries.

# Other Installation Options

- VDT - Virtual Data Toolkit
  - PacMan installer
  - Includes other Grid software
  - <http://www.cs.wisc.edu/vdt/>
- RPM - For Red Hat derived Linux distributions

# Other Sources

- > Condor Manual
- > Condor Web Site
- > condor-users mailing list  
<http://www.cs.wisc.edu/condor/mail-lists/>
- > [condor-admin@cs.wisc.edu](mailto:condor-admin@cs.wisc.edu)

# Publications

- "Condor - A Distributed Job Scheduler," *Beowulf Cluster Computing with Linux*, MIT Press, 2002
- "Condor and the Grid," *Grid Computing: Making the Global Infrastructure a Reality*, John Wiley & Sons, 2003
- These chapters and other publications available online at our web site

# Thank you!

<http://www.cs.wisc.edu/condor>  
[condor-admin@cs.wisc.edu](mailto:condor-admin@cs.wisc.edu)